


Building your Mouseless Development Enviroment



Matthieu Cneude

Contents

Building Your Mouseless Development Environment	10
Who Should Read This Book?	11
What Is a Mouseless Development Environment?	11
What Do You Need to Follow Along?	12
Creating Your Own Cheatsheets	12
Experimenting Is Key	13
Styling Conventions	13
Choose Your Tools	13
In a Nutshell	13
 Part I - Arch Linux	 15
A General Linux Overview	16
Diving Inside Linux	16
The Linux Filesystem	17
Linux Distributions	18
Packages and Repositories	19
Why Arch Linux?	19
The Glory of Rolling Distributions	19
The Arch Linux Community	20
Official Repositories and the Arch User Repositories (AUR)	20
The Fabulous Manual	20
Troubleshooting	21
In a Nutshell	21
Going Deeper	21
 The Power Is In Your Fingers	 22
Efficient Typing: The Two Rules	22
The First Week	23
The Second Week	23
Speed and Accuracy	24
In A Nutshell	24
Going Deeper	24
 Preparing Your System for Arch Linux	 25
Prerequisites	25
Burning the Arch Linux ISO	25
Configuring the Arch Linux Live System	26
Keyboard Layout	27
Connecting to the Internet	28
System Clock	29
BIOS or UEFI?	29
Partitioning the Hard Disk	30
Wiping Your Hard Disk	31
Using fdisk	32

Boot Partition	32
Root and Swap Partition	34
Formatting the Partitions	35
The UEFI Boot partition	35
Root and Swap Filesystems	35
Mounting the Filesystems	36
Mounting the Root Partition	36
UEFI and The Boot Partition	37
Continuing The Installation	37
In a Nutshell	37
Going Deeper	38
Installing Arch Linux	39
Installing the Base Packages	39
Mounting Automatically Partitions With fstab	40
Changing The Root Directory with arch-chroot	40
The Root User's Password	41
Through Time and Space: Configuring the Timezone	41
Choose Your Locale	42
Naming Your New World	43
GRUB, The Linux Bootloader	43
Bootloader With an UEFI	44
Bootloader With a BIOS	44
Enabling the Network Manager	45
Diving in the Shell	45
Command Lines 101	45
Input Output Redirections	46
Pipes	47
Input and Arguments	47
Rebooting The System	48
In a Nutshell	48
Going Deeper	49
Welcome To Arch Linux	50
Connecting to The Internet, Third Round	50
Using a Cable	50
Using the Wi-Fi	51
Installing the Manuals	51
Processes	51
The Init Process: systemd	52
Logging the Init Process With journald	54
Processes As Files	55
Environment Variables	55
Creating A New User	55
The Default Text Editor	57
The Return of the Environment Variable	58
Trying visudo Again	58
First Steps In Neovim	58
Neovim Modes	59
Editing with Neovim	59
Using Sudo	60
Best Practices For Linux Shells	60
Strings	61
Globbing	61
Strings And Quotes	61
In a Nutshell	62
Goind Deeper	62
Links	62

Manual	63
Graphical Interface	64
The X Window System	64
The X What?	64
Installing X	65
URxvt, Our Terminal Emulator	65
Video Terminals and TTYs	66
Installing i3 Window Manager	67
Launching X	67
Installing Fonts	68
Launching i3	68
In a Nutshell	68
Going Deeper	68
 Part II - The Tools	 69
First Steps In Your Mouseless Development Environment	70
Screen Resolution	70
The Basics of i3	71
i3bar and i3status	71
Installing Your Favorite Browser	72
Program Launcher	72
Copy and Paste	72
Pasting From a PDF	73
The Book Companion	73
Troubleshooting	74
In a Nutshell	74
Going Deeper	74
 Configuring URxvt	 75
The Colors of the Terminal	75
Making URxvt Prettier	76
Configuring URxvt	78
Character Fonts	78
Window Default	78
URxvt Daemon	79
In a Nutshell	79
Going Deeper	80
 The Basics of Neovim	 81
Configuring Neovim	81
Neovim Modes	81
Forget the Arrow Keys	82
Switching To Insert Mode	83
Undo And Redo	83
Motions	83
Horizontal Motions	83
Vertical Motions	84
Scrolling commands	84
The Language of Neovim	84
Operators	85
Text Objects	85
A Basic Configuration	85
Neovim's Options	85
Swap Files	86
Undo Tree	86
General Options	86

In a Nutshell	87
Going Deeper	87
Arch Linux Package Managers	88
Operations and Options	88
Official Repositories	88
Updating Your System	89
Removing Packages	89
Searching Packages	90
Pacman's Cache	90
The Arch User Repository (AUR)	90
The Package Manager Yay	91
Clearing Yay Cache	92
Installing Packages With Yay	92
Adding Tabs for URxvt	93
Neovim Language Extensions	93
Pacman Troubleshooting	94
Pacman Configuration	94
In a Nutshell	94
Going Deeper	94
Dotfiles	95
Hard and Symbolic Links	95
Hard Links	95
Symbolic Links	96
The Structure of the Dotfiles Project	96
Our First Bash Script	97
Running A Shell Script	98
Linux Permissions	99
Making Our Dotfiles Public	100
The Protocol SSH	101
Adding Your SSH Public Key to Github	101
In a Nutshell	102
Going Deeper	103
i3: A Deeper Dive	104
How To Use i3?	104
General Organization	104
The Default Shortcuts	105
Configuring i3	106
Configuration Files	106
Default Configuration	106
Program Launcher	107
Moving Windows and Changing Focus	107
Split containers	108
Workspaces	108
Resizing Windows	111
Locking Your Screen	111
Lock, Shutdown, and Reboot Menu	112
Wallpaper	113
Floating windows	113
Colors and Style	114
Scratchpad	115
The i3 Status Bar	116
Managing Your Screen	117
Updating Our Dotfiles	118
In a Nutshell	119
Going Deeper	119

The Z-Shell (Zsh)	120
Framework Or No Framework?	120
Zsh Config Files	121
Basic Zsh Configuration	122
Environment Variables	122
Aliases	123
Options	124
Zsh Completion System	124
Pimp My Zsh Prompt	125
Zsh Directory Stack	126
Zsh, Your New Best Friend	127
Zsh With Vim Flavors	127
Enabling the Vi Mode	127
Changing Cursor	127
Vim Mapping For Completion	128
Editing Command Lines In Neovim	128
Zsh Plugins	128
Zsh Additional Completion	129
Zsh Syntax Highlighting	129
Jumping To A Parent Directory	129
Custom Scripts	129
The Fuzzy Finder fzf	130
Automatically Starting i3	131
In a Nutshell	132
Going Deeper	132
Zsh Documentation	132
Author's dotfiles	132
Improving Your Mouseless Development Environment	133
Improving the Dotfiles Install Script	133
Adding And Configuring Fonts	135
Installing Fonts	135
Changing URxvt's Fonts	136
Changing i3's Fonts	136
Desktop Notifications With Dunst	137
Automatically Mounting Devices	138
Visual Configuration	138
Remapping Your Caps Lock	139
A New Escape Key	139
Two Keys in One	139
The Problem of the External Keyboard	140
Git diff	140
In a Nutshell	141
Going Deeper	141
Neovim: A Deeper Dive	142
Neovim Spatial Organization	142
Buffers	142
Windows	144
Tabs	145
Argument List (arglist)	145
Mapping Keystrokes	146
Jump! Jump! Jump!	148
Jump List	148
Change List	148
Method Jump	148
Repeating Keystrokes	148
Single Repeat	148

Complex Repeat: The Macro	149
Command Line Window	149
Revisiting the Undo Tree	149
Plugins	150
Plugin Manager	150
Closing Buffers Without Closing Windows	152
Managing Windows Easily	152
Navigating Through The Buffer List	152
Manipulating the Undo Tree	152
Automatically Installing the Plugin Manager	152
In a Nutshell	153
Going Deeper	154
The Terminal Multiplexer tmux	155
What's tmux?	155
Why using Tmux?	156
Background Operations	156
More Terminals! Everywhere!	156
Saving tmux Sessions	156
Remote Pair Programming	156
How to use tmux?	157
General Organization	157
tmux Workflow	158
Managing tmux Sessions	158
Configuring tmux	158
The Essential of tmux	159
Increasing The Maximum Output Lines	161
Managing Windows	161
Design	163
Plugins	164
The tmux Plugins Manager	164
Fuzzy Search And Copy with fzf and Extrakto	165
Creating tmux Sessions Automatically	165
i3 Scratchpad Running tmux	166
Choosing Your tmux Session With fzf	167
In a nutshell	168
Going Deeper	168
Neovim Plugins	169
The Language Server Protocol	169
The Plugin coc.vim	169
Extensions to coc.vim	170
tmux completion	170
Fuzzy Finder in Neovim With fzf	170
Navigating files	171
Linter	171
Surrounding	171
Navigating in Open Buffers	172
Text Objects	172
Register History	172
Snippets	172
Search And Replace	173
Status Bar	173
Color Scheme	173
Manual Pages In Neovim	173
The Undo-tree	174
Tmux	174
Startup	174

Git	175
Syntax Highlighting	175
Misc	175
In a Nutshell	175
Going Deeper	176
Mouseless Browsers	177
Lynx	177
Help	178
Navigation	178
History	178
Page Data	178
Bookmarks	178
Downloading	178
Options	179
Quitting Lynx	179
Qutebrowser	179
Basics	179
Navigation	179
Tabs	180
Modes	180
Browser Plugins	180
Firefox	180
Chrome	180
In a Nutshell	181
Going Deeper	181
 Part III - Arch Linux Installer	 182
The System Installer	183
The Project	183
The User Interface	185
Preliminary Configuration	185
Is It the Good Time?	185
Return Code and Operators	185
Output Redirection	187
UEFI or BIOS, That Is The Question	188
Choosing The Hard Disk	189
Bash arrays	190
Pipes	191
Awk	191
Grep	192
Putting Everything Together	192
In a Nutshell	193
Going Deeper	193
 Partitioning and Installing Arch Linux	 194
Size of the Partitions	194
Erasing the Hard Disk	195
Creating Partitions	196
Boot Partition With BIOS or UEFI	196
Automating fdisk	196
Formatting partitions	198
General Case	198
Special Case	198
Generating fstab And Installing Arch Linux	198
The Adventure Continue!	199
The End of the Installer	199

In a Nutshell	200
Going Deeper	200
Creating Users and Passwords	201
Getting Back the Block Devices and the Boot Mode	201
Naming Your New Born System	202
Installing The Bootloader GRUB	202
Clock and Timezone	202
Configuring the Locales	203
Root Password and User Creation	203
Arch Linux Is Now Fully Configured	206
In a Nutshell	207
Going Deeper	207
Installing The Tools	208
The List of Applications	208
Groups of Applications	210
Parsing the CSV	211
Updating the System	212
Installing the Packages	213
Permission For Power: sudo	215
Invoking The Last Installer Script	215
In a Nutshell	215
Going Further	216
The User Installer	217
Usual Linux Directories	217
Installing Packages From the AUR	217
Installing the Dotfiles	219
The One Command Line To Invoke The Installer	219
In a Nutshell	220
Going Deeper	220
This is the End	221

Building Your Mouseless Development Environment

Welcome, Mouseless Developers! Thanks for picking up this book. We'll embark together in a creative journey where we'll build an efficient and mouseless development environment. We'll go from the void of an empty hard disk to a complete system you can mostly use with a keyboard. We'll install everything manually at the beginning, and we'll become an omniscient entity able to summon the whole development environment with one command line at the end.

There are huge benefits to install a complete system by ourselves: we'll learn a tone along the way. It's good to use tools to achieve our goals, but it's even better to know a minimum how they work. It will allow us to modify whatever we want according to our needs.

Knowledge brings flexibility, and flexibility brings efficiency.

Our secret weapon? The command line. We'll use it for almost everything we'll do throughout this book, because it's simply the most powerful tool we can use. It's consistent, it doesn't get out of fashion, it doesn't go in our way, it doesn't dramatically change when new versions are available. It will make your life easier and it will give you tremendous powers. Who doesn't want that?

If you need more arguments to convince you that the command line is what you need to level up to a higher plan of existence (at least), here we go:

1. As developers, we often have to deal with Linux-based systems on servers (or containers) where only the command line is available.
2. Some Command Line Interfaces (CLIs) don't have any graphical interface. To use these programs, we have no choice: we need the shell. Keep the acronym "CLI" in mind: we'll use it often throughout the book.
3. The shell gives you the ultimate power we all seek: automation. It's difficult to automate the movements of your mouse on a graphical interface; it's easier when you deal with plain text, like command lines.

This book will explain everything for you to understand what we're doing, why we're doing it, and how you can personalize your system according to *your* needs. To *your* workflow. To *your* personality. The goal: acquiring the knowledge you need by practicing and experimenting. You'll then be able to transfer this knowledge on whatever development environment you want, even if it's based on another Linux distribution, or if you want to use a more standard IDE instead of Neovim, for example. You can even use most of the tools in this book on macOS.

Now, a bummer: we won't dive deep in everything we'll speak about, or this book would be-

come way too long for all of us to survive it; both the poor writer and the poor readers. Instead, it will focus on showing how the different tools described in this book can work together.

That said, each chapter of the book concludes on a list of resources you can use to dive deeper to your heart's content.

Who Should Read This Book?

Everyone can read this book, but not everybody will get the most value out of it. So, should you read this book?

If you're a beginner in software development, this book is for you. I try to explain everything you need to know to understand what we're doing. Yet, it can be a bit tough on you because it's a lot of information to swallow at once. My advice: go slowly, don't hesitate to try things out, to experiment, to play with the command line. This book will be truly rewarding if you put the effort and some patience.

If you're a seasoned developer but you don't use the command line often, you're at the right place! This book is for you if you want to know more about Linux, the shell, and if a Mouseless Development Environment spark your infinite curiosity. Personally, when I discovered it, the spark looked more like an enlightenment! You think I'm exaggerating? Yes I do! But still.

If you already use the command line intensively and the tools I describe in this book, you might not learn much from it. I'm sorry. It breaks my heart more than yours. That said, Building Your Mouseless Development Environment can help you fill the gaps in your knowledge. Buying the book can be a good way to support my work, too; if you like [my blog](#), my [Github](#), my style, or my limited charisma. You can still offer this book to some poor souls who still work on Windows, too. It's not that I judge you, Windows developers; I was in your situation for decades. It's just that the grass is greener and tastier on the Linux side of the fence.

What Is a Mouseless Development Environment?

The obvious goal of a Mouseless Development Environment is to use less the mouse. It doesn't mean that you can throw your mouses through your windows. First, because you might kill somebody, second because it's not nice to pollute, and third because the mouse is still useful in many cases. I'm not a dogmatic.

Using the mouse is great for some endeavors, like graphical manipulations, video editing, or musical creations. I wouldn't draw in a terminal like I wouldn't write a book with a brush; it wasn't meant for that.

I'm sure you noticed but software engineers deal a lot with text: we spend our time writing code and (hopefully) documentation. It's for this kind of job that our keyboard and our command lines really shine. Text never get out of trend: you can write scripts to parse them and to automate almost whatever you want, thanks to the Holy Shell. Repeat after me: glory to the holy shell!

A Mouseless Development Environment let your hands on the keyboard most of the time, which is a blessing by itself. It's very comfy not to move your hands to reach your mouse, then your keyboard, then your mouse, in an infinite recursion of pain. Even if you think that it doesn't bother you, you'll see the truth by simply trying not to use it. I thought it wasn't

a problem for a long time, but trying to stay on the keyboard definitely changed my way of working.

What Do You Need to Follow Along?

To follow this book, you need to have a *place* to install the whole system and its tools: Arch Linux, URxvt, Neovim, Zsh, tmux, and so on. If you have an empty hard disk, that's a very good container for that. You can use as well a virtual machine if you want to see first how it looks like, or if you just want to follow the book for learning purposes, without the intention to use the final Mouseless Development Environment.

The good new: even if you finish the book on a virtual machine, you'll have a complete installer for the Mouseless Development Environment somewhere on Github. It's something we'll build together, too. You can use it on any standard computer you want, even on a MacBook Pro; tested and approved! If you fall in love with your new shiny environment, as I did, you'll be able to install it quickly on a new computer.

You can create virtual machines using [Virtualbox](#). A virtual machine is a simulated computer using the resources of your physical computer. You can install whatever you want on it, without your physical computer to know about it. Both systems (host and virtual machine) are decoupled as much as possible. In general, virtual machines are great to try different OS or Linux distributions without too much hassle.

If you install the Mouseless Development Environment on a physical computer, I advise you to find a way to be able to read this book while installing everything. You can use another computer, a tablet, or a reader for example.

Creating Your Own Cheatsheets

Since a Mouseless Development Environment focus on using the keyboard when it's appropriate, we'll see many shortcuts (or keystrokes) in this book. We'll go through each tool step by step to understand why and how to use these keystrokes, for you to remember them easily.

Still, you need to be able to come back to these keystrokes if you forget them while using your system. You can of course download already made cheatsheets on the Internet, but in my experience they are not very useful for beginners. You risk ending up in an ocean of keystrokes, not really knowing what are the most important ones, suffocating by too many possibilities.

That's why I would advise you to write your own cheatsheets while going through the book. One for each tool we'll see together. If you need to be convinced, here some arguments:

1. Writing will help you to remember the different keystrokes. You can as well write some comments, categorize them, and even add some personal mnemonics. You can draw something funny near your keystroke. Humour is a great tool to memorize. In short, you'll make these keystrokes *yours*.
2. You can organize them in a way which makes sense to you.
3. When you'll come back to them, you'll feel they are some extension of your brain, not 10923810938 unknown keystrokes downloaded from a random, cold, and sad Internet corner.

I followed this technique when I learnt to build my own Mouseless Development Environment, and I believe that's one of the reason why I never found Neovim or anything else to have a high learning curve. I was then able to get things done in this mouseless system very quickly. Believe me, it's not because I'm a genius; I'm a very standard human.

Similarly, you should as well write somewhere the different command lines you'll see in this book, for the same reasons.

Experimenting Is Key

To learn, you need to practice and experiment. It's a bit more work than watching passively a Youtube video, but it's more effective too. Write your own cheatsheet, comment your configuration files and bash scripts, and don't hesitate to play around with everything. The goal is not to memorize but to understand how it works. Try to modify a command line and see what options you can use, for example.

In two words: be curious.

More you'll learn about the shell in general, easier it will be for you to learn what you can put on top and how it works.

Styling Conventions

There are only a few special styling conventions in this book. You'll know, while reading the book, what to execute in the shell, and what to write in what config files.

Sometimes, `< something >` will pop up in some command lines you need to execute. That is, an expression surrounded with the characters `<` and `>`. For example, `fdisk < your_hard_disk >`, or `su < user_name >`. These are variables you, and only you, know the values of. No worries: I'll tell you each time how you can find the information you need to replace them with the good value.

Choose Your Tools

You can replace whatever tools we'll see in this book with others, but I would suggest you to follow the entire book first to understand how they can work together. Then, you can modify whatever configuration file or replace whatever tool you want because, instead of copying and pasting the configuration of a random person on the Internet, you'll have built your whole system by yourself. You'll have a great control over it.

We'll use Neovim for every editing tasks from the beginning on. If you don't want to learn how it works, you can use Nano instead or whatever else you prefer.

In a Nutshell

We'll begin, in the next chapter, to shape our system to install Arch Linux afterward. What did we learn in this chapter?

- If you're not already using the command line and a Mouseless Development Environment daily, you'll learn many things from this book. Otherwise, you can still try to fill some gaps in your knowledge.
- Create your own cheatsheets to maximize your learning. Make the useful information your own.
- Don't be afraid to experiment with the command lines and the tools we'll see in this book.

Part I - Arch Linux

A General Linux Overview

This chapter is a high level overview of a Linux-based system. Don't worry if you don't remember everything: we'll come back to these concepts in later chapters.

Let's first see some generalities about the system we'll build. More precisely we'll answer these questions:

- What's Linux?
- What's a Linux distribution?
- What are repositories?
- What's the shell?

I know. It might look very boring for seasoned Linux users. Who knows? Maybe you'll learn two or three things from it. Of course, you're free to skip this chapter if you want to.

Diving Inside Linux

Linux is an Operating System (OS) like Windows or macOS. The heart of any OS is a program called the *kernel*. It's the interface between you (and the applications you use) and the hardware of your computer. Thanks to the kernel, the programs running on your system can communicate with the hard disk or the memory, for example.

For you to speak with the kernel around a good cup of tea, you need another interface: the *shell*. For example, you can create a file using the shell, which is in fact you, the user, asking politely to the kernel to create something on the hard disk. The kernel might accept your request or deny it ruthlessly.

The shell is an interpreter, which means that it can read *command lines* and perform actions depending on them, like running a program or creating a file. You can write these command lines yourself using a standard input. The most obvious standard input is a keyboard, but you can use a script too, a file containing command lines to execute. We'll come back to that later in this book.

There are many shells out there you can use, like Bash or Zsh. We'll use Zsh to install Arch Linux, and then we'll continue with Bash for a while.

To use the shell, guess what? We need another interface. In computing, we heavily rely on layers of abstractions, and we use interfaces to communicate between these layers. This time, a *terminal emulator*, or *terminal*, will allow you to access the shell. Again, we'll come back to that later, and explain more in detail what it is.

You'll see throughout the book and in many other places the sentences "execute in a shell",

“execute in a terminal”, “run in a shell”, or “run in a terminal”. These sentences are all synonym. It means that you need to:

1. Type the command the book provide in a terminal.
2. Press the `ENTER` key to interpret the command.

Then, our friend the shell will return an *exit code*, not directly visible, and possibly an *output* which will be displayed in the terminal by default.

The shell should be your friend. She should be your best buddy, because she can do a lot for you. Yes, the shell is a feminine character in my world, but your shell can be whatever you want. The sky's the limit.




Figure 1: Kernel, shell, and terminal

The Linux Filesystem

The directories in a Linux-based system are all children of a special directory called *root directory*. In the filesystem, it will seem hidden at first, known only under the name `/`. Don't be fooled! It's the most important directory because it contains every other one. A directory to rule them all!

Each direct children of the root directory have precise purposes we'll see all along the book.

How does look a path to a file (or filepath) in a Linux based system? For example: `/etc/zsh`. The first `/` is our powerful root directory, supporting everything else as Atlas support the world. The directory `etc` is a direct child of the root directory, and the directory `zsh` is itself

a child of `etc` . Every `/` which are not the first character of the filepath is a *directory separator*. It's here to indicate a new level of directory. Confusing? I told you the root directory tried to hide!

If you run in a terminal the command `cd /etc` , you'll move from the directory you were to the directory `/etc` . It will become your *working directory*, called as well *current directory*. Each time you move from directory to directory, your working directory change too. You can think of it as the place you are, and you can jump from place to place. To display what your current directory is, you can run the command `pwd` (for `p`rint `w`orking `d`irectory) in a terminal.

There are many special directories in a Linux based system. Here are two you'll encounter very often:

- `.` - Represent your working directory.
- `..` - Represent the direct parent of your working directory.

If you run `cd .` , you'll jump from the directory you're in to the directory you're in. Was it useless? Absolutely! Now, if you run `cd ..` , you'll move to the parent directory of the working directory. Similarly, if you run in a shell `cp ./cat_picture ./file` , it means that you'll `c`o`p`y the file named `cat_picture` in your working directory and "move" it to your working directory. It means that you'll have two identical files in your working directory with different names.

Linux Distributions

If you install an OS like Windows or macOS, it will include installing the kernel, a shell, and many other programs as well. For example, the Windows installer will provide a desktop environment, a bundle of programs sharing the same graphical interface. It includes the desktop where you can place shortcuts of your programs, often a status bar with a truckload of stuff in it, a wallpaper, and all of these things. The OS doesn't let you choose what programs and what desktop you want to install.

A Linux distribution, commonly called *distro*, is a Linux kernel, a shell, and many programs too. The difference: there are 200+ Linux distros you can choose from. For example, the distro Ubuntu will install some version of the Linux kernel, the shell Bash, a desktop environment called Unity, and many other programs, like some Amazon ads. Who doesn't want Amazon ads in a development environment? Everybody? Thanks, Ubuntu.

Sometimes, when people use the word "Linux", it's often not very clear what they are speaking about: a Linux distro? The kernel? The Linux philosophy? What's important to remember is this: there are not many differences between Linux distributions. You can roughly summarize them as follows:

- The set of applications a Linux distro install on your system.
- The set of applications available in the repositories.
- The philosophy attached to the distro, especially the way you can update the Linux kernel, the shell, and all the programs installed.
- The package manager used to install packages.

Some distros are build on top of other ones. For example, Ubuntu is build on top of another distro called Debian. It means that both will share many similarities.

Arch Linux is a distribution too, but it only installs a very minimal Linux-based system. You'll have the kernel, the shell, a set of programs commonly used in the shell, and a package manager. No desktop environment, no Amazon thingy, not even a graphical environment to display Graphical User Interfaces (GUIs). We'll install all of that ourselves, and at the same time we'll see how these layers work together. By understanding this, you'll be able to install whatever you want and personalize everything following your craziest desires.

Packages and Repositories

We were speaking about programs and applications. In the Linux world you'll see as well often the term *package*. A package is a compressed archive bundling every file for a given application. This archive has metadata to know how to install the application itself. These packages are stored in some locations called *repositories*. To download and install packages, we'll use a *package manager*.

Different Linux distros often use different repositories and package managers. For example, Ubuntu uses APT as a package manager, and Arch Linux uses Pacman (for [Package Manager](#)).

Why Arch Linux?

Why will we install Arch Linux in this book, and not another distro?

First, Arch Linux users have often the reputation to be arrogant persons, pointing the finger to everybody who doesn't use Arch Linux. They always have a great pleasure to precise that they use Arch Linux in any discussion, even if it's about the price of the cheese at your favorite supermarket. I wanted to invite you in this very closed club. My pleasure!

On top of the fact that installing Arch Linux can teach a great deal about Linux systems in general, using it in our daily work has many advantages, too.

The Glory of Rolling Distributions

As we saw above, the different Linux distros out there have different ways to manage everything installed on your system, from the Linux kernel to the packages installed.

Arch Linux is a rolling distribution; the packages are kept up-to-date as much as possible. You can be pretty confident that you'll have access to the latest versions of your favorite applications. The small downside: you'll have to update your system pretty often, every week or every other week. Trust me, that's a small price to pay for having the most up-to-date system possible.

If you listen to the urban legend that Arch Linux is "very unstable" because of that, don't believe it. Some friends run Arch Linux for years (otherwise they wouldn't be my friends of course!) and I do, too. We never had any problems doing so. I used Windows from Windows 98 to Windows 7, macOS, and some other Linux distributions like Ubuntu. My conclusion: Arch Linux is the most stable system I've ever had.

The Arch Linux Community

The Arch Linux community is great. Arch Linux folks have solid knowledge in many areas, and they have very useful tips too. The best place to fall in admiration for the Arch Linux community is the [Arch Linux wiki](#). It's simply the best resource you'll find regarding Linux, even if you use another distro.

Official Repositories and the Arch User Repositories (AUR)

The official repositories of Arch Linux propose many packages. I'm not exaggerating: you'll find most of the applications you need in there. Even if you don't, you'll have access to the Arch User Repository too, an unofficial repository where you'll find everything and anything. No need to install, compile, and update manually the obscure applications you use. An AUR helper can do that for you in one command line.

The Fabulous Manual

Almost every CLI out there has at least a `man` ual page describing what it is, how it works, and what it can do. This documentation is often very complete: I advise you to look at it as soon as you need some explanations.

Here's how to access the manual:

```
man <command>
```

For example, you can run `man cd` in a shell to read the manual page of the command `cd`. If you're a beginner, it might be a bit difficult to read at first. When you'll understand better some concepts distilled in this book, you'll find `man` less cryptic.

Keep in mind: you don't have to read everything. Often, reading the description is enough to understand what a CLI can do for you. If you're searching precise *option*, `man` can be very helpful too. A command line option is often a single letter prefixed by a minus `-`, or a word prefixed by a double minus `--`. Adding options will modify the behavior of a command.

A command line has often some kind of option to display some concise help as well, often something like `--help` or `-h`. The output is often shorter and more condensed than the man page for the same command. To understand what I mean, try to run:

```
ls --help
```

Manuals are often divided into sections. To look at a precise section, you can run the following:

```
man <section> <command>
```

I won't go into more details here. If you want to know more about `man` (including this idea of section), you should read the `man` page of `man` by running in a terminal:

```
man man
```

It's inception directly in your shell!

Troubleshooting

Installing manually everything following a book can lead to some problems. Heck, using a computer lead to some problems. In my experience, using an automatic installer for a Linux distro can lead to some problems, too. In short: we can have problems.

If you run into unexpected errors and weird behaviors, you should take a look at these resources:

1. The [Arch Wiki](#) is, as we just saw, the best resource for almost everything, except the meaning of life. If you search on the Internet “ Arch Linux”, you’ll end up very often in the Arch wiki.
2. The [official Arch Linux website](#) is the place to go when you have problems updating your system. It describes the manual interventions you need to do from time to time. Everything is always explained clearly.
3. The [Arch Linux forum](#) is a good place to go if you have any question. Be sure that your problem is related to Arch Linux and that you’ve done some research before posting there.

In a Nutshell

What did we learn in this chapter?

- The kernel is the heart of an Operating System (OS). You can communicate with the kernel using the shell, which will read and interprets commands you type in a terminal emulator.
- In the Linux filesystem, every directory are children of the root directory, named `/`. A filepath looks like this: `/etc/kernel`, where the first `/` is the root directory and the subsequent `/` are directory separators.
- A package is an installer for an application. They are located in repositories you can access using the Internet. You need a package manager to install the application contained in a package.
- Linux distros are not that different from each others. Mostly, they have a different philosophy to handle packages, different package managers, and different repositories. They come bundled with different applications, too.

The next chapter will be a short one: it will teach you how to have good typing techniques. It sounds boring I know, but if you really want to be efficient with your keyboard, this is the way to go.

Going Deeper

- [Arch Wiki](#)
- [Arch Wiki - Arch Linux](#)
- [Arch Linux Website](#)
- [Arch Linux Forum](#)
- [Manual pages for Arch Linux packages](#)

The Power Is In Your Fingers

You remember Frodo? He saved Middle Earth by throwing a ring into a volcano. He was quite small, but how important!

This chapter is quite small too, and, like Frodo, very important.

You're reading right now a book to build a Mouseless Development Environment, so it's pretty obvious I'll advocate to use your keyboard as much as you can. The goal is simple: you should move your hands as little as possible from this magical device.

Knowing how to type properly is a stepping stone to master a keyboard driven workflow. You can practice the following techniques while reading the book. When you begin to replace your typing techniques, stick to it as much as you can. Don't come back to your old (and bad) habits.

The benefits are huge. Massive. Colossal!

- You'll feel way more comfy while developing the next trendy application which will make you rich. What a secret pleasure to have our hands on the keyboard!
- The room for progression (speed and accuracy) will increase drastically.

If you look at pianists playing, you'll notice that they never look at their hands. They don't want to spend brain power into this kind of details; their hands follow their will as fast as possible.

The analogy holds very well for us, computer addicts. When you are in a state of flow, focused on your craft, you don't want to be interrupted every five minutes by moving your hand to the mouse or by looking at your fingers. You don't want to think about all of that. You want to create, nothing else.

Let's see how we can achieve that.

Efficient Typing: The Two Rules

It's satisfying to see your typing techniques improving day after days, months after months, and even years after years. These techniques are easy to learn but difficult to master.

The first rule you need to be aware of: placing your hands correctly. The keys `a`, `s`, `d`, `f` and `j`, `k`, `l`, `;` are called the *row keys*. They are the starting points for your hands. From there, you'll be able to grab any other key as efficiently as possible.

You'll notice that there are little bumps on the keys `f` and `j` on your keyboard: they are indicators for you to know where to put your indexes. When they are at the good position, simply place the other fingers on the other row keys.




Figure 2: Correct hand placements - [colorblind version](#)

The second rule you need to train for: try not looking at your keyboard while you're typing. Of course, if you don't remember where a key is, look at it, but only after trying blindly where you think it is. We want to train your muscle memory, here.

I was only typing with two fingers before trying to follow these two rules. It felt really weird to use these new techniques at first; now, I wouldn't type differently. It's efficient, it's comfortable, it's great!

The First Week

When you decide to use the two rules we saw above, you need to try to follow them *all the time*. We need 100% commitment here. If you surprise yourself using your bad techniques again, which will happen, don't worry: simply come back to the good ones. This is part of the learning process, not a horrible failure cursing your whole family on five generations.

The first three days are the most difficult. You'll alternate between good and bad technique without even noticing it. You'll do mistakes. You'll be slower. That's great! It's how you'll learn.

Fortunately, at the end of the week, the amount of mistake you'll make will decrease, and the need to watch the keyboard will slowly disappear.

The Second Week

You'll notice during the second week the amount of mistakes decreasing even more, and you won't dare looking at your keyboard while typing. At the end of the week, you'll see your typing

speed already improving. The good feelings of reward will begin to please your brain. That's what we all want.

Speed and Accuracy

During your two weeks of initial training, you shouldn't focus on speed or accuracy. Just type, as much as you can, and don't worry about anything else yet. Not even about the mistakes you're making.

Only then, when you'll feel comfortable enough, you can shift your focus on speed and accuracy: how fast you can type while making as few mistakes as possible.

In A Nutshell

What did we learn in this chapter?

- We need to focus on our work, not on our hands. Learning the good typing techniques will help you tremendously in this regard.
- The first rule: place your hands correctly on your keyboard, using the row keys.
- The second rule: don't look at your hands while typing. Even if you don't remember where a key is, don't look at it before trying to get it without looking.
- Only train for speed and accuracy when you're used to these new typing techniques.

In the next chapter, we'll begin to prepare our hard disk for Arch Linux.

Going Deeper

As always, to learn as fast as possible, you need to practice. This book ask you to type many command lines, so you'll have many occasions to learn these good typing techniques.

You can as well use typing software to have concrete data about your speed and accuracy. Here are my favorites:

- [Type Racer](#)
- [Online Typing Test WPM](#)
- [Speed Coder](#)

Preparing Your System for Arch Linux

It's time to get our hands dirty! In this chapter, we'll prepare our system for Arch Linux to feel at home. More precisely, we'll answer these questions:

- How to burn an Arch Linux ISO on a USB key?
- How to configure the Arch Linux live system?
- How to partition a hard disk using `fdisk`?
- How to create and mount new filesystems?

This is the very beginning of the journey, so I'll try to explain most things in details. Ready for the challenge? Let's go!

Prerequisites

To install our new Mouseless Development Environment, we'll need the following:

1. A computer with a 64-bit CPU (not older than 10 years).
2. An empty hard disk, or a hard disk you'll wipe out (at least 20 GB to feel comfy).
3. A USB key (at least 1 GB).
4. A second computer, a tablet, or whatever device allowing you to read this book while installing everything.
5. Internet access on your second device is highly recommended, in case you have a problem in the first chapters.

As we spoke about already, the computer can be a virtual machine or a physical one.

Burning the Arch Linux ISO

If you wonder what the heck an ISO is, it's a file representing a physical disk. We can burn ISO files on a real disk, which means copying everything from the ISO to the disk.

To install Arch Linux, we need to start (or *boot*) our computer using the Arch Linux live system, and install the OS from there. Here's how to do that:

1. Download the [Arch Linux ISO](#). Find your country, click on the first link, and download `archlinux -< the_last_release_date >-x86_64.iso`.
2. Verify that the ISO is the official one, and has not been modified by horrible hackers. You need to generate the MD5 hash of the ISO and compare it with the one provided on the download page.

- On Windows, you can use the command `CertUtil -hashfile < path_to_ISO_file > MD5`.
 - On Linux, you can use the command line `md5sum < path_to_ISO_file >` to get the MD5 of the file.
3. Burn the ISO on a USB key. You can use:
 - On Windows: [rufus](#).
 - On Ubuntu: Startup Disk Creator.
 - On another Linux distro: [UNetbootin](#).
 4. Change the boot order to read USB devices before your hard disk. You can do that by using your BIOS interface.
 5. Plug your USB key and start your computer.

For the fourth point, you can normally access the BIOS interface by hitting a precise key when your computer starts. If you're lucky, it might be displayed at startup long enough for you to see it. It's often `ESC`, `DEL` or `F10`. Every computer is different on this point, so I can't really help you more.

When you found a way to access your BIOS' interface, search any option concerning the boot. You can normally change the boot order of the devices: the goal is to put your USB device before your hard disk. When you're done with that, there'll be some options to save your changes and restart your computer.

If you did everything correctly, your computer will boot on your USB key, and you'll see a menu inviting you to install Arch Linux. Let's do it!

Configuring the Arch Linux Live System

Your screen will then spit out many green "OK", telling you what systemd is starting. We'll speak about systemd later in this book. Then, you'll end up in a shell prompt. It will look like this:

```
root@archiso ~ #
```

A prompt is a set of character indicating that you can execute command lines. Here, it indicates what user you are (`root`) and the hostname `archiso`. The hostname is the name of your system as it will appear on a network.

Welcome to the Arch Linux live system! What's a live system, you might wonder? For now, nothing has been installed on your hard disk; instead, the system you have access to is running from your computer's memory (RAM). Many distros propose a live system for you to test it even before installing anything.

It means as well that everything you're doing in the live system won't survive if you shut down or restart your computer, except if you write deliberately something on your hard disk. For example, any program you'll install on the live system won't be installed on your hard disk.

The shell we're using now is called the "Z shell", or Zsh. It's similar to the most famous shell, Bash. We'll dig more into Zsh later in this book, too. Here are some functionalities you can use with the shell:

- You can enter any command in a shell prompt and execute it by hitting `ENTER`.
- The shell can save a command line history:

- You can quickly access executed command lines with the keys `ARROW UP` and `ARROW DOWN`.
- You can search through the history with `CTRL+r`.
- You can use Zsh auto completion by hitting `TAB`.
- To stop a running command line, you can use the keystroke `CTRL+c`. Be careful with it: patience is sometimes safer. Think about what your command is doing before interrupting it, and what's the price you might have to pay by stopping it before it's done.

You are now the *root user*. You can create different users on a Linux system, but the root user is special. It has almost all powers a user can have. Don't confuse the root user and the root directory `/` we saw in the previous chapter. A filesystem and a bunch of users are different things, even if they have both root as name.

By being the root user, you're in fact the demigod (or demigoddess!) of the "archiso" live system. Why only demi? Because a user can't control directly the kernel. The kernel is the real master around here.

Keyboard Layout

By default, you'll end up with the American keyboard layout on the Arch Linux live system. If you're not comfortable with it, you can change it. Here are the two commands you can use to do so:

1. `ls /usr/share/kbd/keymaps/**/*.map.gz` - List all layout available.
2. `ls /usr/share/kbd/keymaps/**/*.map.gz | grep "< your_language_code >"` - Filter all layout with grep.

For example, if I want to use a French keyboard layout, I can do this:

```
ls /usr/share/kbd/keymaps/**/*.map.gz | grep "fr"
```

The result will look like the following:

```
/usr/share/kbd/keymaps/i386/azerty/fr-latin1.map.gz
/usr/share/kbd/keymaps/i386/azerty/fr-latin9.map.gz
/usr/share/kbd/keymaps/i386/azerty/fr.map.gz
/usr/share/kbd/keymaps/i386/azerty/fr-pc.map.gz
/usr/share/kbd/keymaps/i386/bepo/fr-bepo-latin9.map.gz
/usr/share/kbd/keymaps/i386/bepo/fr-bepo.map.gz
/usr/share/kbd/keymaps/i386/dvorak/dvorak-ca-fr.map.gz
/usr/share/kbd/keymaps/i386/dvorak/dvorak-fr.map.gz
/usr/share/kbd/keymaps/i386/qwertz/fr_CH-latin1.map.gz
/usr/share/kbd/keymaps/i386/qwertz/fr_CH.map.gz
/usr/share/kbd/keymaps/mac/all/mac-fr_CH-latin1.map.gz
/usr/share/kbd/keymaps/mac/all/mac-fr.map.gz
/usr/share/kbd/keymaps/sun/sunt5-fr-latin1.map.gz
```

You can choose then your layout by using the filename without the file extension `map.gz` :

```
loadkeys <filename>
```

For example, for my French keyboard layout, I would run the command `loadkeys fr- latin1`.

Connecting to the Internet

What would we be without Internet? Nothing more than unconscious amoebas lost in a chain of misconceptions we call reality. Without Internet, I wouldn't even be able to write the previous sentence!

To connect to this infinite amount of knowledge and cat pictures, we need some network device. The command `ip-link` will show you the truth about them. Let's try to run in the terminal:

```
ip link
```

Something like this will be output:

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode
   DEFAULT group default qlen 1000
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: enp0s25: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc fq_codel
   state DOWN mode DEFAULT group default qlen 1000
   link/ether f0:de:f1:84:f3:a6 brd ff:ff:ff:ff:ff:ff
3: wlp3s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP
   mode DORMANT group default qlen 1000
   link/ether 08:11:96:0a:12:b4 brd ff:ff:ff:ff:ff:ff
```

Don't worry about the first result `lo`. The second one `enp0s25` is my network device for Ethernet cable; it often begins with an 'e'. The third one `wlp3s0` is for the Wi-Fi; this one often begins with a "w". We don't have the same computer, so it's more than likely we won't have the same network devices either.

If you want to connect to Internet using a cable, simply plug it. For the Wi-Fi, you can run the following:

```
iwctl
```

A new shiny prompt will appear. It works similarly to the shell prompt you had before. Here are the commands you can use:

1. `device list` - List all your Wi-Fi devices.
2. `station <device> get-networks` - Replace `<device>` with the name of one of yours, to get a list of networks you can connect to with this particular device.
3. `station <device> connect "<network>"` - Replace `<device>` and `<network>` by the device you want to use with the network you want to connect to.
4. If a password is required, enter it.
5. `station <device> show` - Show you if you're indeed connected.
6. `exit` - Exit `iwctl`.

If it failed, you might get the friendly message `Operation failed`. Don't worry and try again: even demigods fail, sometimes. Demigoddesses too. Another reason why we're only "demi" here, and not the full version.

Let's now verify if you're really connected. Millions of people use this command all around the world as we speak, to verify their Internet connection:

```
ping -c 5 thevaluable.dev`
```

It will send 5 requests to the address `thevaluable .dev` and display the following:

```
64 bytes from v22017064675050008.supersrv.de (185.183.156.38): icmp_seq=1
ttl=58 time=32.6 ms
64 bytes from v22017064675050008.supersrv.de (185.183.156.38): icmp_seq=2
ttl=58 time=32.6 ms
64 bytes from v22017064675050008.supersrv.de (185.183.156.38): icmp_seq=3
ttl=58 time=32.1 ms
64 bytes from v22017064675050008.supersrv.de (185.183.156.38): icmp_seq=4
ttl=58 time=32.9 ms
64 bytes from v22017064675050008.supersrv.de (185.183.156.38): icmp_seq=5
ttl=58 time=32.2 ms
```

If you've got something similar, it means you're connected to the infinite Internet! Lucky you.

If it doesn't work, try:

```
ping -c 5 185.183.156.38
```

If this last command works, it means that you can't resolve addresses like `thevaluable .dev` to its IP address `185.183.156.38` . In that case, try to run a dhcp client:

```
dhcpcd &
```

What's `thevaluable .dev` ? It's my blog, to show you how to do some subtle product placement in your own book.

System Clock

Let's now use one of the oldest protocol on the Internet. Anxious? Don't worry, Internet is the only stuff which never breaks in the computing world. Not like enterprise Java code.

Run this command:

```
timedatectl set-ntp true
```

We synchronized our system clock with the Network Time Protocol. That's all. Next!

BIOS or UEFI?

You remember the interface you used to change the boot order of your devices, to boot your USB key before another device?

This was the interface of a *firmware*, a piece of software tightly linked to a piece of hardware. In this precise case, this firmware is called a BIOS (for **B**asic **I**nput **O**utput **S**ystem) and it's directly embedded in a chip in your motherboard. The BIOS is the first software running when you're booting (starting) your computer.

This BIOS verify that your hardware works properly. It makes available an interface as well for you to configure some basics, like the order of the booting devices. It will as well run the *bootloader*, another program which is used to boot an operating system, like Arch Linux.

That said, the BIOS is now considered deprecated for a better alternative, called *UEFI* (for **U**nified **E**xtensible **F**irmware **I**nterface). BIOS and UEFI are two different firmwares, but confusion arises when the UEFI pretends to be a BIOS, sometimes using the term “BIOS” on its own interface! It’s often an ironic attempt not to confuse people who are used to have a BIOS.

Long story short, because the firmware of your motherboard start the bootloader of Arch Linux, we need to know if you have a BIOS or an UEFI on your computer. To verify that, simply run in the shell:

```
ls /sys/firmware/efi/efivars
```

If you see the error `no such file or directory`, it means that your computer has a BIOS. If it output a bunch of files, you have an UEFI. You need to remember what kind of firmware you have to install Arch Linux properly. My advice: write it somewhere, on a piece of paper for example. We’ll need this information later to create our boot partition and install the boot-loader itself.

We’ve just learn the first piece of software which run on our computer. You know what we should do? Create a diagram we’ll complete as we go, to describe the boot process of a Linux-based system:




Figure 3: The boot process: BIOS and UEFI

This process is followed by all Linux distros out there. The green box under the arrow indicate the program (or group of programs) used during this process, sometimes specific to Arch Linux.

Partitioning the Hard Disk

Enough explanations. Let’s do some serious stuff now. First, let’s run the following in the shell:

```
lsblk
```

This command will display your *block devices*. A block device is a file which represent a physical device. Confused? How can a device, such as a hard disk, be represented as a file?

A Linux-based system tries hard to have a consistent way for you, the applications running on your system, the kernel, and the hardware, to interact with each other. In other words, Linux

offers a consistent *interface* between you and your physical system, and between the programs running and your physical computer.

Indeed, to facilitate this deep and intimate relationships, there is an important principle to understand on a Linux system: **everything is a file**. You don't need to ask yourself how a program read or write a device, like your hard disk: the answer will always be via a file.

For example, you can try to run the following in your shell:

```
cat /etc/proc/stat
```

The output will give you the CPU status at the precise time you read the file. That's what I mean when I speak about an interface between you (or some running programs) and a device, here your CPU.

We'll come back to this idea over and over: it's central to any Linux-based systems. For now, if you look at the output of our command `lsblk`, you'll see something similar to that:

NAME	MAJ:MIN	RM	SIZE	RO	TYPE	MOUNTPOINT
sda	8:0	0	465.8G	0	disk	
sda1	8:1	0	200M	0	part	/boot
sda2	8:2	0	16G	0	part	[SWAP]
sda3	8:3	0	389.6G	0	part	/

- `sda` is a file representing my hard disk (`TYPE` `disk`).
- `sda1` to `sda4` are partitions of the hard disk (`TYPE` `part` for partition). A partition is a virtual division of a physical hard drive.

If your hard disk is empty, you shouldn't see any partition.

You can see the files representing your hard disk in the `/dev` directory ("dev" for "device").

Wiping Your Hard Disk

If your hard disk is not empty, you need to delete everything on it. All its data will be lost, forever. It can take a long time depending on the size and the type (mechanical or SSD) of the hard disk. As an example, for a mechanical disk of 1 GB, it took me 4 hours to entirely wipe it.

If you're good with that, run the following in your shell:

```
dd if=/dev/zero of=<your_hard_disk> status=progress
```

You need to replace `< your_hard_disk >` with yours. For example, considering the output above when I run the command `lsblk`, the name of my hard disk is `/dev/sda`. As a result, to wipe it, I need to run `dd if=/dev/zero of=/dev/sda status=progress`.

If you have multiple hard disks, be sure to select the good one! The command `dd` is merciless; its sweet little nickname is "Disk Destroyer". It won't ask you anything and wipe out whatever you give. You don't want to mess up with `dd`!

The command copy everything from the file `/dev/zero` to whatever file representing your hard disk (for me, `/dev/sda`). The file `/dev/zero` is a file containing an infinity of zeros; new zeros will be created each time you try to read from it. In our precise example, your hard

disk will be filled with zeros till it's full. At that point, the command `dd` will stop and your disk will be "clean".

You can use as well `/dev/urandom` instead of `/dev/zero` : it will fill your hard disk with random numbers instead of zeros. It takes even more time, so use this solution only if you don't want anybody to be able to restore the removed data of your hard disk.

Using fdisk

It's time to virtually break our hard disk in pieces! We need to partition it for Arch Linux to come and make its nest. To do so, we'll use a simple and powerful Command Line Interface (CLI) called `fdisk` .

The first time I installed Arch Linux, `fdisk` felt very mysterious to me. It took me a long time to partition my disk, and I couldn't repeat the process afterward without going through the same pain. The horrible truth is: `fdisk` is, in fact, quite easy to use.

First, we need a *partition table*. This table is read by the OS to know what are the partitions on the disk. Then, we'll create three partitions:

1. A *boot partition*, where the bootloader (to start the OS) will be installed.
2. A *swap partition*, a place on your hard disk used as memory when your RAM is full. It's not perfect since the RAM is quicker than a hard disk, but it's better than nothing.
3. A *root partition*, where everything else will be stored: Arch Linux files, your cat pictures, your secret poetry, and so on.

Don't confuse the root partition with the root directory or the root user we saw earlier. These concepts are different.

Let's now run the following command in your terminal:

```
fdisk <device_name>
```

For example, I can see that my disk is called `sda` when I run `lsblk` , so I should run the command `fdisk /dev/sda` . From there, a new prompt will appear. It will look like this:

```
Command (m for help):
```

You can execute special one-letter commands in `fdisk`. Nothing will be written on your hard disk when doing so, except if you use the command `w` (for `w`rite). So if you do a mistake, don't panic: you can abort everything with `CTRL + c` if you didn't write anything yet.

Let's create now a GPT partition table. Let's execute the following command in `fdisk`'s prompt:

```
g
```

It can be surprising to use one-letter commands, but it's how `fdisk` works.

Boot Partition

Now that we've chosen our partition table, let's create the boot partition. Enter the following command:

```
n
```


At that point, `fdisk` will ask you the partition number. The default is normally `1`, so just hit `ENTER`.

As command lines go, `fdisk` is very curious. It loves asking questions. It will ask you then the first sector of the disk where the partition should begin. By default, it's the beginning of the disk, and that's what we want: simply hit `ENTER` again.

The last question will ask you what should be the last sectors for the partition. It's not obvious, but you can enter the size you want for your partition here. For the boot partition, 512 MB is enough, so let's type:

```
+512M
```

You'll come back to `fdisk`'s prompt after that. The whole operation should look like this, including the creation of the partition table:

```
Command (m for help): g
Created a new GPT disklabel (GUID: <some_GUID>).

Command (m for help): n
Partition number (1-128, default 1):
First sector (2048-71567838, default 2048):
Last sector, +sectors or +size{K,M,G,T,P} (2048-71567838, default 71567838)
: +512M

Created a new partition 1 of type 'Linux filesystem' and of size 512MiB.

Command (m for help):
```

We need now to change the type of the boot partition, depending on your boot firmware (UEFI or BIOS). To do so, enter the following command:

```
t
```

To list all the partition types possible, you can use the command `L`. To quit the list, use `q`.

UEFI Boot Mode

If you have an UEFI, we want the boot partition to be of type `EFI System`. It's normally the first one in the list, so you should enter:

```
1
```

Here's how the whole operation should look like:

```
Command (m for help): t
Selected partition 1
Partition type (type L to list all types): 1
Changed type of partition 'Linux filesystem' to 'EFI System'
```

BIOS Boot Mode

If you have a BIOS, we want the boot partition to be of type `BIOS boot`. It's normally the fourth in the list, so you should enter:

```
4
```

Here's how the whole operation should look like:

```
Command (m for help): t
Selected partition 1
Partition type (type L to list all types): 4
Changed type of partition 'Linux filesystem' to 'BIOS boot'
```

Root and Swap Partition

We have now configured the boot partition. Don't forget: we still didn't write anything on the hard disk, so if you quit `fdisk` now, you'll need to do everything again from the beginning.

You should still see `fdisk`'s prompt at that point. We can now create a SWAP partition. If you have more than 8 GB of RAM and you don't do anything which takes a lot of memory (like video editing, for example), you can skip the creation of a SWAP partition entirely. That being said, I always like to create one of 8 GB, because "you never know".

Let's create a new partition with the command:

```
n
```

When asking the partition number (normally `2` by default), just hit `ENTER`.

When `fdisk` asks you the first sector of the partition on the disk, the default is just after the end of the last partition. This is again what we want, so simply press `ENTER`.

Then, we can enter the size of the partition. It needs to be bigger than 512 MB. I want 8 GB, so I enter:

```
+8G
```

When it's done, you're back to the `fdisk`'s prompt. Let's move forward by creating the root partition. Execute the following:

```
n
```

Then, it's even easier: hit `ENTER` for every question `fdisk` asks. For the size of the partition, it will take by default whatever is left on the disk. When you're back at the prompt, enter the following command:

```
p
```

You should see something like that:

Device	Start	End	Sectors	Size	Type
/dev/sda1	2048	1050623	1048576	512M	BIOS boot
/dev/sda2	1050624	17827839	1677726	8G	Linux filesystem
/dev/sda3	17827840	195352513	193569729	923G	Linux filesystem

If you have an UEFI, the only difference will be the type of the first partition: it will be `EFI System` instead of `BIOS boot`.

Again, the device name, start, end, and sectors columns can have different values on your system. The things we have in common: the number of partition created, their sizes (except for the last one), and their types.

It's time to write our changes on the hard disk. Enter the following command to write the partitions:

```
w
```

That's it! You can now quit `fdisk` using the command `q`. When you're back in the shell prompt, you can run the command `lsblk` again to see the newly created partitions.

Formatting the Partitions

The partition table and the partitions themselves are now created, but it's not enough for Arch Linux to use them. Each partition needs to be formatted with a specific filesystem.

Basically, a filesystem group the data on a disk in files and directories. Using the filesystem, the OS can then display the content of a disk as we're used to. Different OS support different filesystems; for example, Windows support primarily the `NTFS` filesystem.

I won't describe every filesystem available here, it would be too cumbersome and quite boring. We'll simply format our partitions using the filesystems we want, thanks to the CLI `mkfs` (`mk` for `make`, `fs` for `file system`).

The UEFI Boot partition

If you have an UEFI, we need first to format the boot partition using the filesystem `FAT32`. Let's run in the shell:

```
mkfs.fat -F32 <your_boot_partition>
```

For example, my boot partition is `/dev/sda1`, so I run: `mkfs.fat -F32 /dev/sda1`.

Root and Swap Filesystems

If you have a BIOS, no need to create a filesystem for your boot partition. It will be done automatically later. For the Swap partition, you need to run the following commands whatever firmware you have, BIOS or UEFI:

```
mkswap <your_swap_partition>
swapon <your_swap_partition>
```

For example, my swap partition is `/dev/sda2`, so I replace `< your_swap_partition >` from the above commands with `/dev/sda2`.

Let's create next an `Ext4` filesystem for the root partition with the command:

```
mkfs.ext4 <your_root_partition>
```

Your root partition is normally the last partition listed when you run `lsblk`. It's as well the biggest partition you have.

This operation can take some time. You don't have to input anything, just let `mkfs` running. Our partitions are now formatted! Glory and joy!

Mounting the Filesystems

If you run `lsblk` again, you'll see an empty `MOUNTPOINT` column. What's that?

For example, if I run `lsblk` on my current system, where Arch Linux is already running, I get this:

NAME	MAJ:MIN	RM	SIZE	RO	TYPE	MOUNTPOINT
sda	8:0	0	465.8G	0	disk	
sda1	8:1	0	200M	0	part	/boot
sda2	8:2	0	16G	0	part	[SWAP]
sda3	8:3	0	389.6G	0	part	/

A mountpoint is a directory where we can access the filesystem of a precise partition. For example, you can see that the mountpoint of `/dev/sda1` is `/boot`. It means that:

- If I list (read) the files in the directory `/boot`, the files on the partition `/dev/sda1` will be listed.
- If I create (write) a file in the directory `/boot`, the files on the partition `/dev/sda1` will be written.

We've created partitions but we can't access them yet. To do so, we need to mount our partitions on precise directories.

Mounting the Root Partition

First, let's mount our root partition `/mnt`. Your root partition is normally the last partition of your disk, as well as the biggest. Let's run the following in the shell:

```
mount <your_root_partition> /mnt
```

If you run the command `ls /mnt`, you'll see that your partition is totally empty, except for a `lost+found` directory. A Linux based system uses this directory on every partition to recover bits of corrupted data if needed.

UEFI and The Boot Partition

If you have an UEFI, you need to mount the boot partition on a new directory. Let's first create this directory:

```
mkdir -p /mnt/boot/efi
```

`mkdir` stands for “make (`mk`) a directory (`dir`)”. What's the option `-p`, you might wonder? Try to run `mkdir --help` in your shell. You'll see the argument `-p` and a brief explanation: `no error if existing, make parent directories as needed`.

The CLI `mkdir` only creates the directory `efi` by default. If the directory `boot` doesn't exist, it returns an error. The option `-p` automatically create the directory `efi` and its `p`arent if it doesn't exist. It doesn't return any error if any directory exists already.

Let's now mount our boot partition on our newly created directory:

```
mount <your_boot_partition> /mnt/boot/efi
```

You can run `lsblk` to see if your mountpoints are correct.

Continuing The Installation

We've now built a comfy nest for Arch Linux to come and change your life.

If you want to continue your installation later and stop your computer, you'll have to go through the part **Mounting the Filesystems** again, after loading the Arch Linux live system from your USB key and connect to the internet using `iwctl`. For now, we don't have any way to automatically mount the partitions or connect to Internet when the computer is starting.

It's not that big of a deal. If your brain is about to explode, make yourself a good cup of tea, do some Yoga, go into the forest singing with the squirrels, and come back when you're ready to install the one and unique Arch Linux.

In a Nutshell

We'll install in the next chapter the basic packages we need: the Linux Kernel and some basic tools. We learned, in this chapter:

- A prompt invite you to run command lines.
- An ISO is a file which represents a whole disk (CD-ROM, hard drive...).
- The Arch Linux live system (like any Linux live system) doesn't run from your hard disk but from your RAM.
- A BIOS is a firmware on your motherboard. Modern motherboards implement a different firmware called UEFI.
- You can display information about your network devices using the CLI `ip`.
- You can make requests on a network with the CLI `ping`. Very handy to test your Internet connection!
- You can copy any content to any destination with the command `dd`.
- You can manage your partitions using `fdisk`.

- You can create new filesystems using `mkfs` .
- You need to mount your partitions on directories to access them, using the command `mount` .

Going Deeper

Taking the habit to read the manual of the CLIs you use will bring you an insane amount of knowledge regarding Linux-based systems. For example, you can read the description of `fdisk` by running `man fdisk` .

Don't worry if you don't understand everything. The most important: take the habits to look at the manual. You'll understand more and more what you're reading there as you go through this book. You don't have to read everything; often, only the description is enough at first. Don't put too much pressure on yourself.

Part II - The Tools

Configuring URxvt

Let's begin to configure the different tools we use. One of the most important is our terminal emulator URxvt. Let's configure it first, to make it more pleasing to the eyes. After all, it should motivate you to do some good work, not to switch off your computer as fast as you can, your stomach rotating on itself by such ugliness.

We'll see in this chapter:

- How to change the colors of our terminal.
- How to configure the basics for URxvt.
- How to use URxvt as a daemon.

You won't recognize your terminal emulator after that!

The Colors of the Terminal

We can configure URxvt using a general purpose configuration file for X. First, let's run in a terminal:

```
mkdir ~/.config/X11
```

We'll put every configuration files for X our new directory `X11`. Let's now create the configuration file itself:

```
touch "$HOME/.config/X11/.Xresources"
```

The file `.Xresources` will contain configurations for applications directly client of X, like URxvt. Don't forget the dot `.` in the filename! This file is called a "dotfile"; you can see it as a synonym of "configuration file" in the Linux world. Be careful: not every configuration file has a filename beginning with a dot!

Many Linux users put all their configuration files online, in projects they often call "dotfiles". We'll do that in the **subsequent chapters**, too. For now, let's go in the folder we just created:

```
cd $HOME/.config/X11/
```

If you try to run `ls` in your terminal, nothing will be displayed. Why? We created a file in there! Every file beginning with a period `.` are hidden files in Linux. You need to do `ls -a` (`a` ll) to display them.

Let's now edit the configuration files. Run in the shell:


```
nvim ~/.config/X11/.Xresources
```

Let's add some colors for the foreground, background and the cursor, by writing these lines in the file `~/.config/X11/.Xresources`

```
*.foreground:    #C6C6C6
*.background:    #1C1C1C
*.cursorColor:   #444444
```

The symbols `*` means that the configuration will be applied to every application using this file. To reload the file `.Xresources` and apply the changes, run the following in another terminal:

```
xrdb -merge $HOME/.config/X11/.Xresources
```

What does it mean? I'm glad you ask! The option `-merge` will merge the configuration you specified in `.Xresources` with the local settings of the applications using this configuration. To see the whole final settings set, you can run `xrdb -query -all`.

You need to restart URxvt itself to display the changes you've just made.

It would be nice to merge our customized `.Xresources` file at startup. Let's edit again the file `~/.xinitrc`:

```
nvim ~/.xinitrc
```

Let's add this line *before* `exec i3`:

```
xrdb -merge $HOME/.config/X11/.Xresources
```

Why do we need this line before `exec i3`? We want `i3` to:

1. Replace the shell process which is started automatically after login with an `i3` process.
2. Run continuously `i3` in the background.

To do so, we always need the line `exec i3` to be at the end of the file!

Making URxvt Prettier

URxvt is a terminal supporting 256 colors. It might be not enough for you (some terminal support True Color, or 16 millions of them), but for me it's already too much. So many choices!

Let's configure every other color you can configure for a terminal emulator. Many other CLIs will use these colors, too. For example, here are mine you can add to `.Xresources`:

```
! black 236 238
*.color0: #303030
*.color8: #444444

! red 167
*.color1: #d75f5f
*.color9: #d75f5f

! green 108
*.color2: #87AF87
*.color10: #87AF87

! yellow 221
*.color3: #ffd75f
*.color11: #ffd75f

! blue 110
*.color4: #87afd7
*.color12: #87afd7

! magenta 146
*.color5: #afafd7
*.color13: #afafd7

! cyan 75
*.color6: #afd7ff
*.color14: #afd7ff

! white 239 15
*.color7: #e4e4e4
*.color15: #ffffff
```

You can copy this configuration directly from [the book companion](#).

If you don't like these colors, you make me very sad. That said, you can go on the fantastic website called [terminal sexy](#) to help you configure the colors of your terminal easily. You can even import and export your [Xresources](#) file back and forth.

The exclamation marks `!` are used for inserting comments. Everything after `!` won't be interpreted. I've added some information using comments for each color.

The first color on the line just after each comment is the darkest version of the color, and the line just below it the lightest version. It means that if an application wants to display light magenta, it will use the `color13` we've configured.

Regarding the comments themselves, the numbers just after the color names are called "Xterm color codes" (xterm is the standard terminal emulator for X). It's a number between 1 and 256 which represents a color. You might need to use this color code for application which doesn't support hexadecimal colors (`#ffffff` for example). It's useful to have a consistent color scheme across your system.

You can find [all the colors you can use here](#).

Configuring URxvt

Enough colors! If you didn't notice already, the actual font used by URxvt is not really readable in some situations. Let's change that.

Character Fonts

You can display all fonts installed on your system with the command:

```
fc-list
```

You can filter the list using `grep` as we saw in previous chapters, or you can navigate through it using a pager like `less`. If you use `grep`, you can use the option `-i` to filter without case sensitivity. It means that you can use uppercase and lowercase characters without `grep` making any distinction between them.

Now, let's add the following line in the file `$HOME/.config/X11/.Xresources` :

```
URxvt*font: xft:DejaVuSansMono:size=14:antialias=true
```

Reload the configuration with `xrdb` as we saw above and restart URxvt. Your font normally changed. You can see here that we've written `URxvt*` at the beginning of the line, which means that the following configuration is only applied to URxvt.

If the spaces between the letters are too wide, you can try to add the following line:

```
URxvt*letterSpace: -1
```

Window Default

Let's now improve URxvt's window itself by adding these three lines in our `.Xresources` :

```
URxvt*borderLess: false
URxvt*externalBorder: 0
URxvt*internalBorder: 4

URxvt*scrollBar: false
```

With this configuration we first hide the external border of the window. Then, we set an internal border to create some nice padding. These borders are transparent by default, but you can add a color with the option `URxvt*borderColor` .

We hide the `scrollBar`, too: it's quite ugly, and we'll see other way to scroll up and down.

Let's display as well more lines in the terminal by adding this line:

```
URxvt*saveLines: 5000
```

It's now possible to scroll through 5000 lines of output in our terminal!

Let's see next how to make our terminal even faster: that's partly why URxvt is such a good choice.

URxvt Daemon

URxvt is a very light terminal emulator. You can even run it as a daemon (program running in the background) as soon as you log in with your user.

We can use instances of the client `urxvtc` to connect to the daemon `urxvtd` (the server). It means that if the daemon crash, every client process connected to it will crash, too. It might look like a big problem, but it's not:

- We'll install tmux later, allowing you to persist tmux sessions even if you close the terminal emulator itself.
- It never happened to me in years of use.

We should run the daemon `urxvtd` at the beginning of our X session. To do so, let's modify again our file `.xinitrc` :

```
nvim ~/.xinitrc
```

Next, let's add this line at the beginning of the file:

```
urxvtd -o -q -f
```

This line will run URxvt's daemon each time we launch our X server with these options:

- `-q` for `quiet` - By default, `urxvtd` display a welcome message. Even if it's nice and welcoming, I see it as useless noise.
- `-o` for `opendisplay` - Keep the daemon process running at all time.
- `-f` for `fork` - Bound URxvt to its control socket.

If you log out, it's important to note that the X server is reseted and `urxvtd` killed. As a result, if you modify some configuration for URxvt and it doesn't seem to work, sometimes you only need to log out and log in again to see the change. It happened to me pretty often when changing the fonts in `.Xresources` , for example.

In a Nutshell

We'll dive more into Neovim in the next chapter. In this chapter, we saw:

- How to configure the colors for our terminal and every program using `.Xresources` as a configuration file (dotfile).
- How to list all the fonts install on our system with `fc-list` .
- How to change the font for URxvt.
- How to modify slightly URxvt's window.
- What's the URxvt daemon and how it works.

Going Deeper

- [Arch Wiki - rxvt-unicode](#)
- [Arch Wiki - X resources](#)